

# VE370 RC (Week 12)

Li Shi

2020.11.27

# Before we start ...

## Recitation Class On-site Evaluation

- Nov 19 – Practical Teaching Workshop



- Nov 27 – RC On-site Evaluation (today)

# Content

- Key concept 1: Motivation of virtual memory (VM)
- Key concept 2: Page table (PT)
- Key concept 3: Translation look-aside buffer (TLB)

# Key concept 1: Motivation of VM

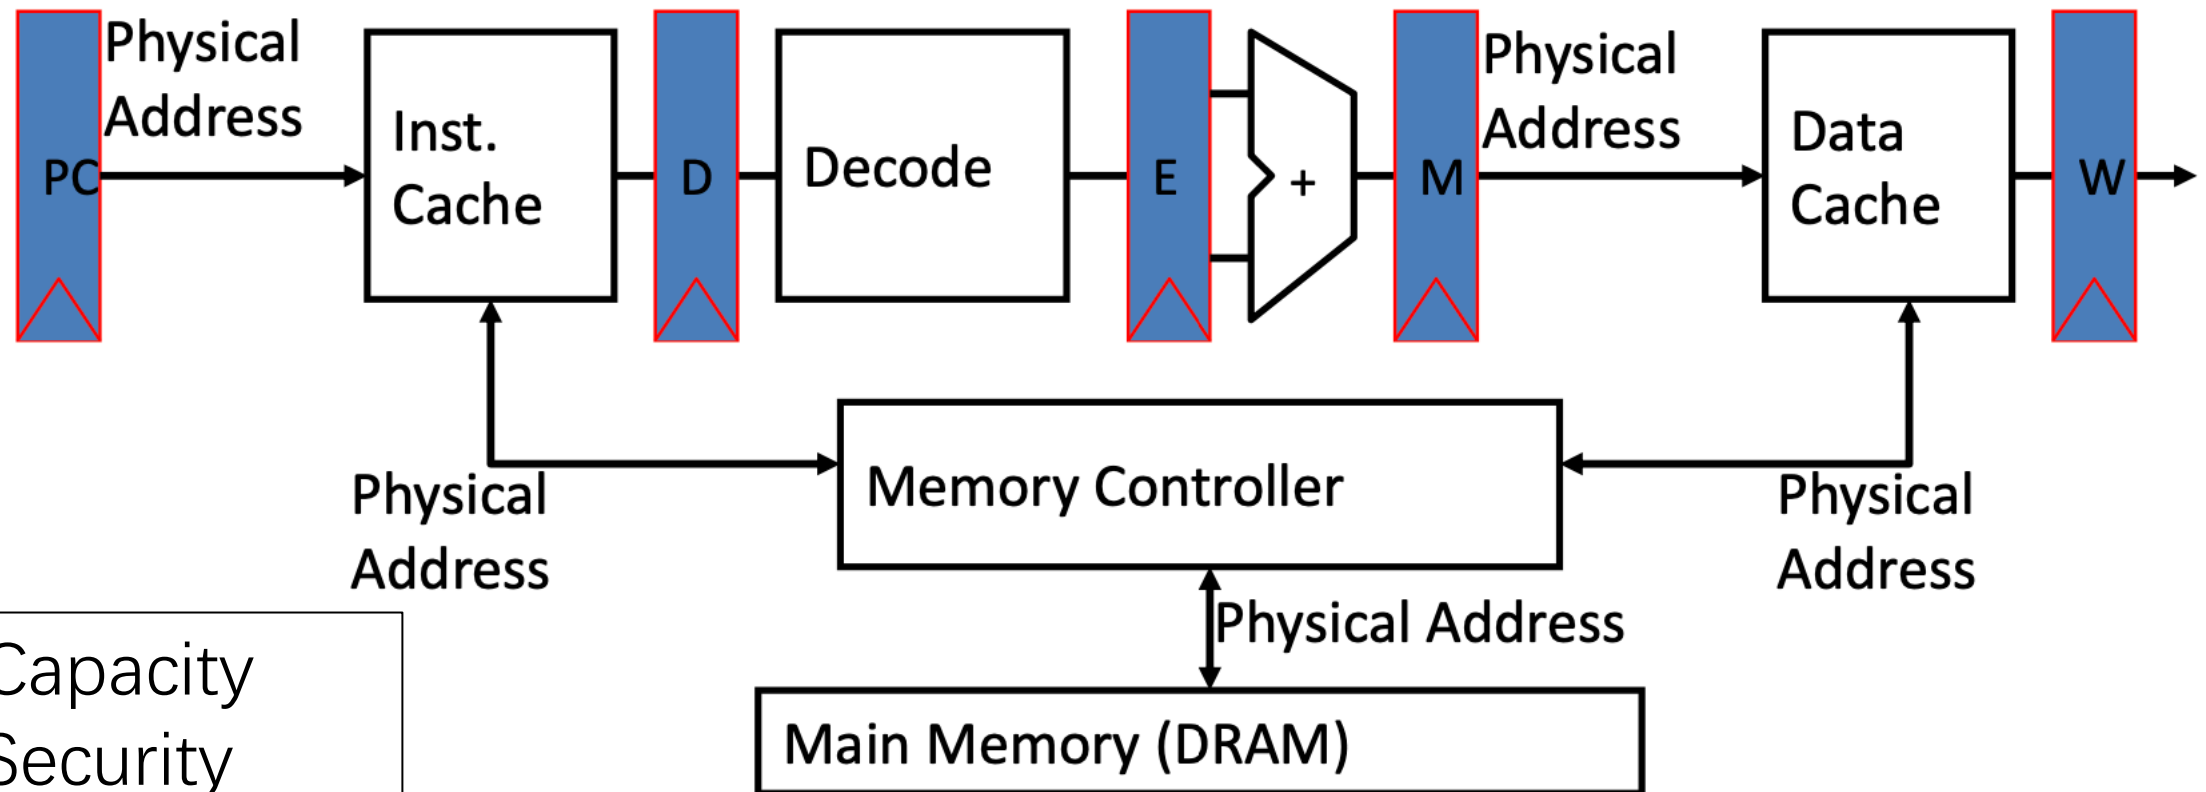
## Issues with memory

- Issue 1: Capacity
  - Large program on hard drive vs. small capacity of memory
- Issue 2: Security
  - Program A may access to memory possessed by program B
- Issue 3: Complexity
  - We need to manage memory quickly and efficiently, but CPU already has many other issues

# Key concept 1: Motivation of VM

## Issues with memory from a historical view

- Solution 1: Bare machine



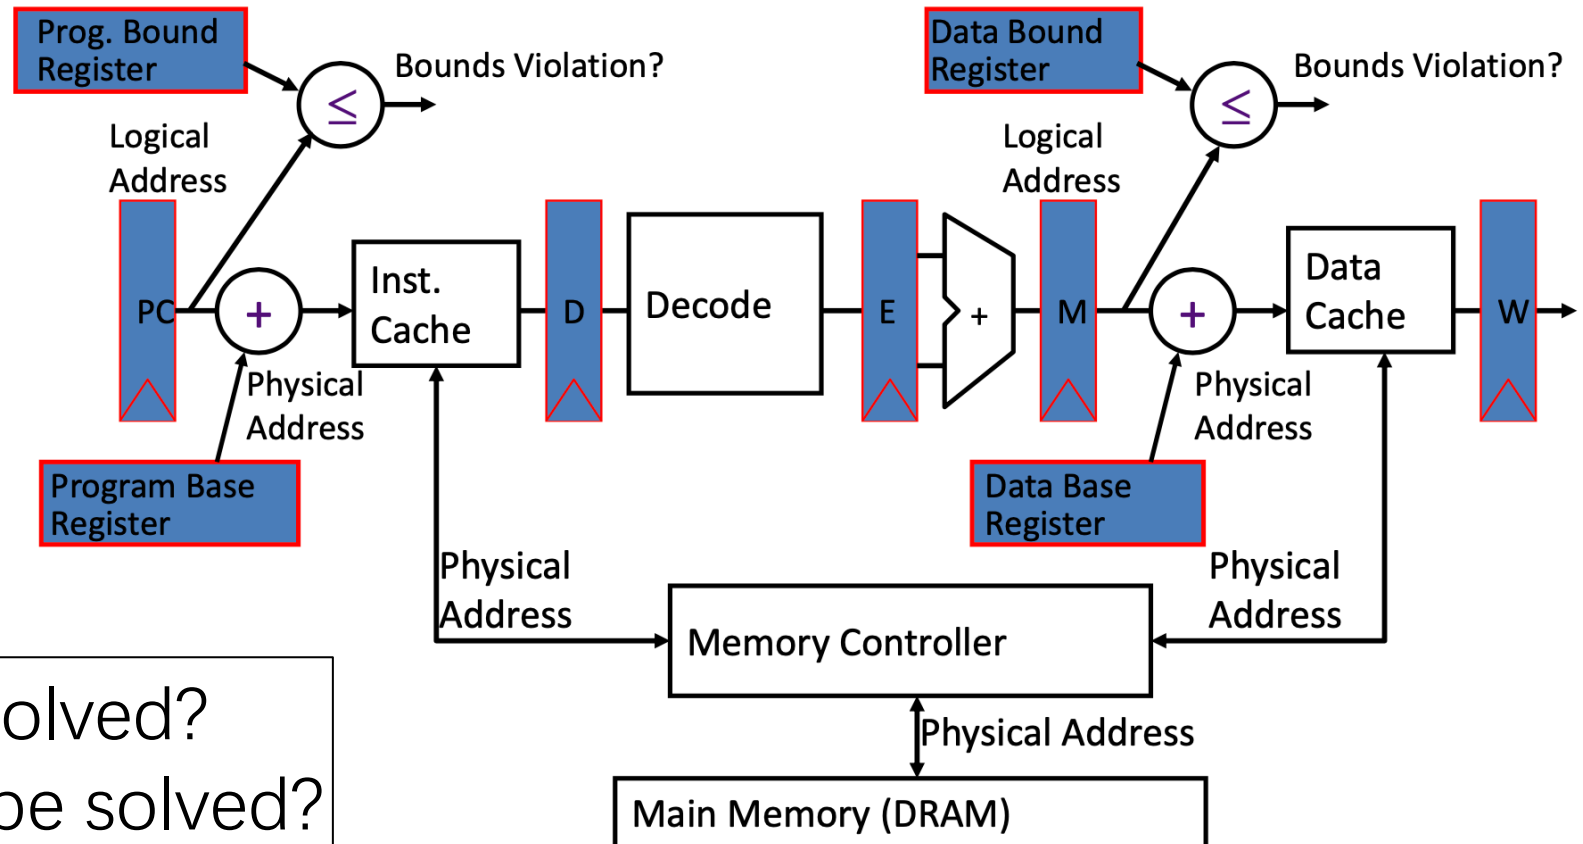
- 1) Capacity
- 2) Security
- 3) Complexity

This slide will not be covered in the exam.

# Key concept 1: Motivation of VM

## Issues with memory from a historical view

- Solution 2: Base and bound machine




What issues **can** be solved?  
What issues **cannot** be solved?

# Key concept 1: Motivation of VM

## Issues with memory from a historical view

- Ultimate solution: Virtual Memory with Page Table



What issues **can** be solved?  
What issues **cannot** be solved?

# Key concept 2: Page Table

## Game: quick ask and quick answer

- Q1: Where is page table located at?  
A. Register file      B. Cache      C. Main memory      D. Disk
- Q2: Who possesses a page table?  
A. A program      B. A process      C. The cache      D. The CPU
- Q3: What is the typical size of one page?  
A. 16 Byte      B. 4 KB      C. 512 KB      D. 4 MB
- Q4: Usually what kind of associativity is used?  
A. Directed map      B. 2-way      C. 4-way      D. Full
- Q5: When a page fault occurs, where is the possible destination to find the page?  
A. Swap space      B. OS kernel      C. Main memory      D. TLB



# Key concept 2: Page Table

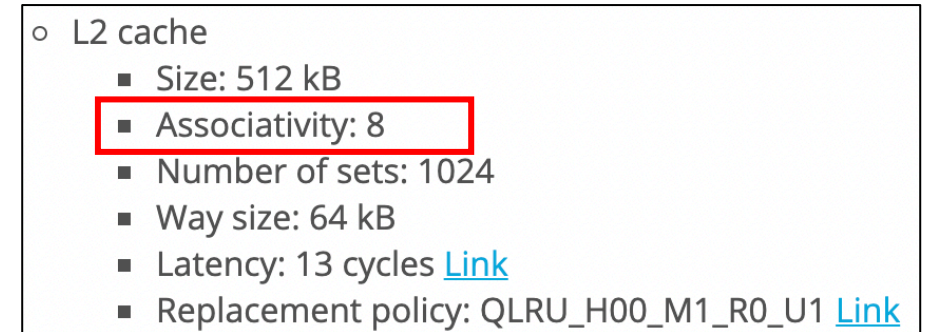
## Game: quick ask and quick answer

- Q1: Where is page table located at?  
A. Register file      B. Cache      C. Main memory      D. Disk
- Q2: Who possesses a page table?  
A. A program      B. A process      C. The cache      D. The CPU
- Q3: What is the typical size of one page?  
A. 16 Byte      B. 4 KB      C. 512 KB      D. 4 MB
- Q4: Usually what kind of associativity is used?  
A. Directed map      B. 2-way      C. 4-way      D. Full
- Q5: When a page fault occurs, where is the possible destination to find the page?  
A. Swap space      B. OS kernel      C. Main memory      D. TLB

# Key concept 2: Page Table

## Beyond these "simple" questions ...

- Q4: Usually what kind of associativity is used? **D. Full (associativity)**
- Compare with CPU cache (e.g., L2 cache of Intel Core i5-1035G1, shown in the following figure).



- Why not use full associativity in cache?
  - Reason: Must search all entries to find a hit – Reduce performance
- Why do we use full associativity in page table?
  - Goal: Reduce page fault rate  $\times$  penalty time (very large)

# Key concept 2: Page Table

## Beyond these "simple" questions ...

- Q3: What is the typical size of one page? **B. 4 KB**
- But ... why 4 KB? How about a larger size? (Homework 8.3.3)

- Why we need to make it large?
  - T14 – Page 18

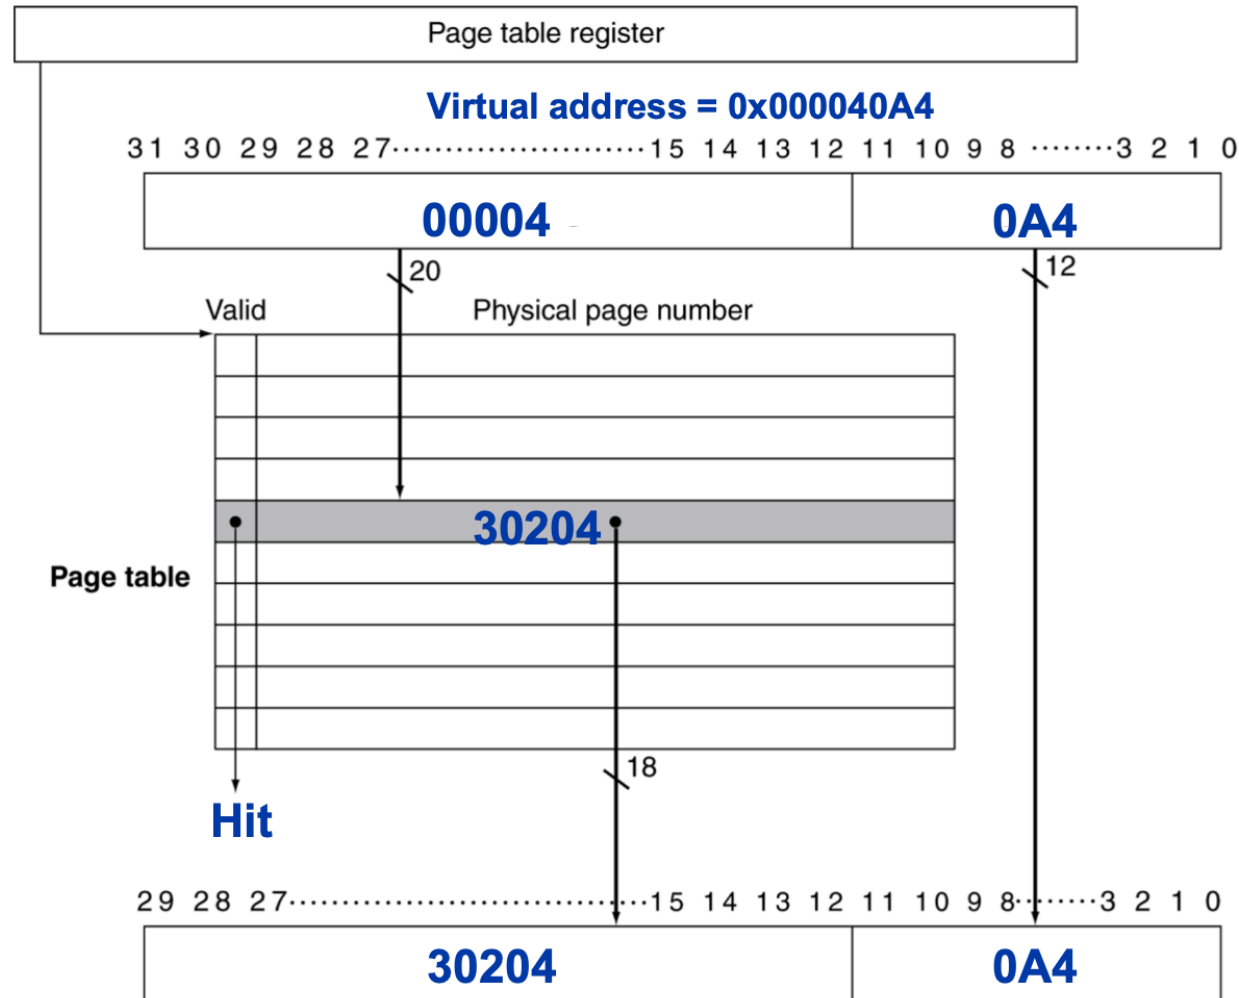
Most of the time is for getting the first word in the page – access time – very long

- Should have **large page size**, so one access fetches more data, also reduces page fault rate

- Why not too large?

- Internal fragmentation (Not all memory in page is used)
- Larger page fault penalty (more time to read from disk)

# Key concept 2: Page Table Address translation



# Key concept 2: Page Table

## Class exercise 1 – Simulate a page table

- Page table: 4KB page size, LRU replacement, stored in the first page in physical memory.
- 16KB physical memory (14-bit physical address)
- 1MB virtual memory (20-bit virtual address)
- How many entries are there in this page table?
- Suppose initially we have the page table on the next slide, simulate the page table when we access the following virtual addresses.

0x00F0C, 0x01F0C (w), 0x02F0C,  
0x00100, 0xFF000 (w), 0x01FFF

# Key concept 2: Page Table

## Class exercise 1 – Simulate a page table

0x00F0C, 0x01F0C (w), 0x02F0C,  
0x00100, 0xFF000 (w), 0x01FFF

Page Table

# (*)	V	D	Physical page number
0	1	0	1
1	1	0	2
2	0	–	Disk
...	0	–	Disk
255	0	–	Disk

Physical Memory

# (*)	Content
0	Page table
1	...
2	...
3	...

Note (\*): Index field does not exist in the real page table and physical memory.

# Key concept 2: Page Table

## Class exercise 1 – Simulate a page table

0x00F0C, 0x01F0C (w), 0x02F0C,  
0x00100, 0xFF000 (w), 0x01FFF

Page Table

#	V	D	Physical page number
0	1	0	1
1	1	0	2
2	0	–	Disk
...	0	–	Disk
255	0	–	Disk

Physical Memory

#	Content
0	Page table
1	...
2	...
3	...

Physical address: 0x1F0C

# Key concept 2: Page Table

## Class exercise 1 – Simulate a page table

0x00F0C, 0x01F0C (w), 0x02F0C,  
0x00100, 0xFF000 (w), 0x01FFF

Page Table

#	V	D	Physical page number
0	1	0	1
1	1	1	2
2	0	-	Disk
...	0	-	Disk
255	0	-	Disk

Physical Memory

#	Content
0	Page table
1	...
2	...
3	...

Physical address: 0x2F0C



# Key concept 2: Page Table

## Class exercise 1 – Simulate a page table

0x00F0C, 0x01F0C (w), 0x02F0C,  
0x00100, 0xFF000 (w), 0x01FFF

Page Fault

Page Table

#	V	D	Physical page number
0	1	0	1
1	1	1	2
2	1	0	3
...	0	-	Disk
255	0	-	Disk

Physical Memory

#	Content
0	Page table
1	...
2	...
3	...

Physical address: 0x3F0C

# Key concept 2: Page Table

## Class exercise 1 – Simulate a page table

0x00F0C, 0x01F0C (w), 0x02F0C,  
0x00100, 0xFF000 (w), 0x01FFF

Page Table

#	V	D	Physical page number
0	1	0	1
1	1	1	2
2	1	0	3
...	0	-	Disk
255	0	-	Disk

Physical Memory

#	Content
0	Page table
1	...
2	...
3	...

Physical address: 0x1100

# Key concept 2: Page Table

## Class exercise 1 – Simulate a page table

0x00F0C, 0x01F0C (w), 0x02F0C,  
0x00100, **0xFF000 (w)**, 0x01FFF

Page Fault

Page Table

#	V	D	Physical page number
0	1	0	1
<b>1</b>	<b>0</b>	<b>0</b>	<b>Disk</b>
2	1	0	3
...	0	0	Disk
<b>255</b>	<b>1</b>	<b>1</b>	<b>2</b>

Physical Memory

#	Content
0	Page table
1	...
<b>2</b>	<b>...</b>
3	...

Physical address: 0x2000

# Key concept 2: Page Table

## Class exercise 1 – Simulate a page table

0x00F0C, 0x01F0C (w), 0x02F0C,  
0x00100, 0xFF000 (w), 0x01FFF

Page Fault

Page Table

#	V	D	Physical page number
0	1	0	1
1	1	0	3
2	0	0	Disk
...	0	0	Disk
255	0	1	2

Physical Memory

#	Content
0	Page table
1	...
2	...
3	...

Physical address: 0x3FFF

# Key concept 3: TLB

Game: quick ask and quick answer

- Full name of TLB?

- Matching:


- |               |   |                        |
|---------------|---|------------------------|
| • L1 cache    | ● | ● Cache of cache       |
| • L2 cache    | ● | ● Cache of page table  |
| • Main memory | ● | ● Cache of main memory |
| • TLB         | ● | ● Cache of disk        |

# Key concept 3: TLB

Game: quick ask and quick answer

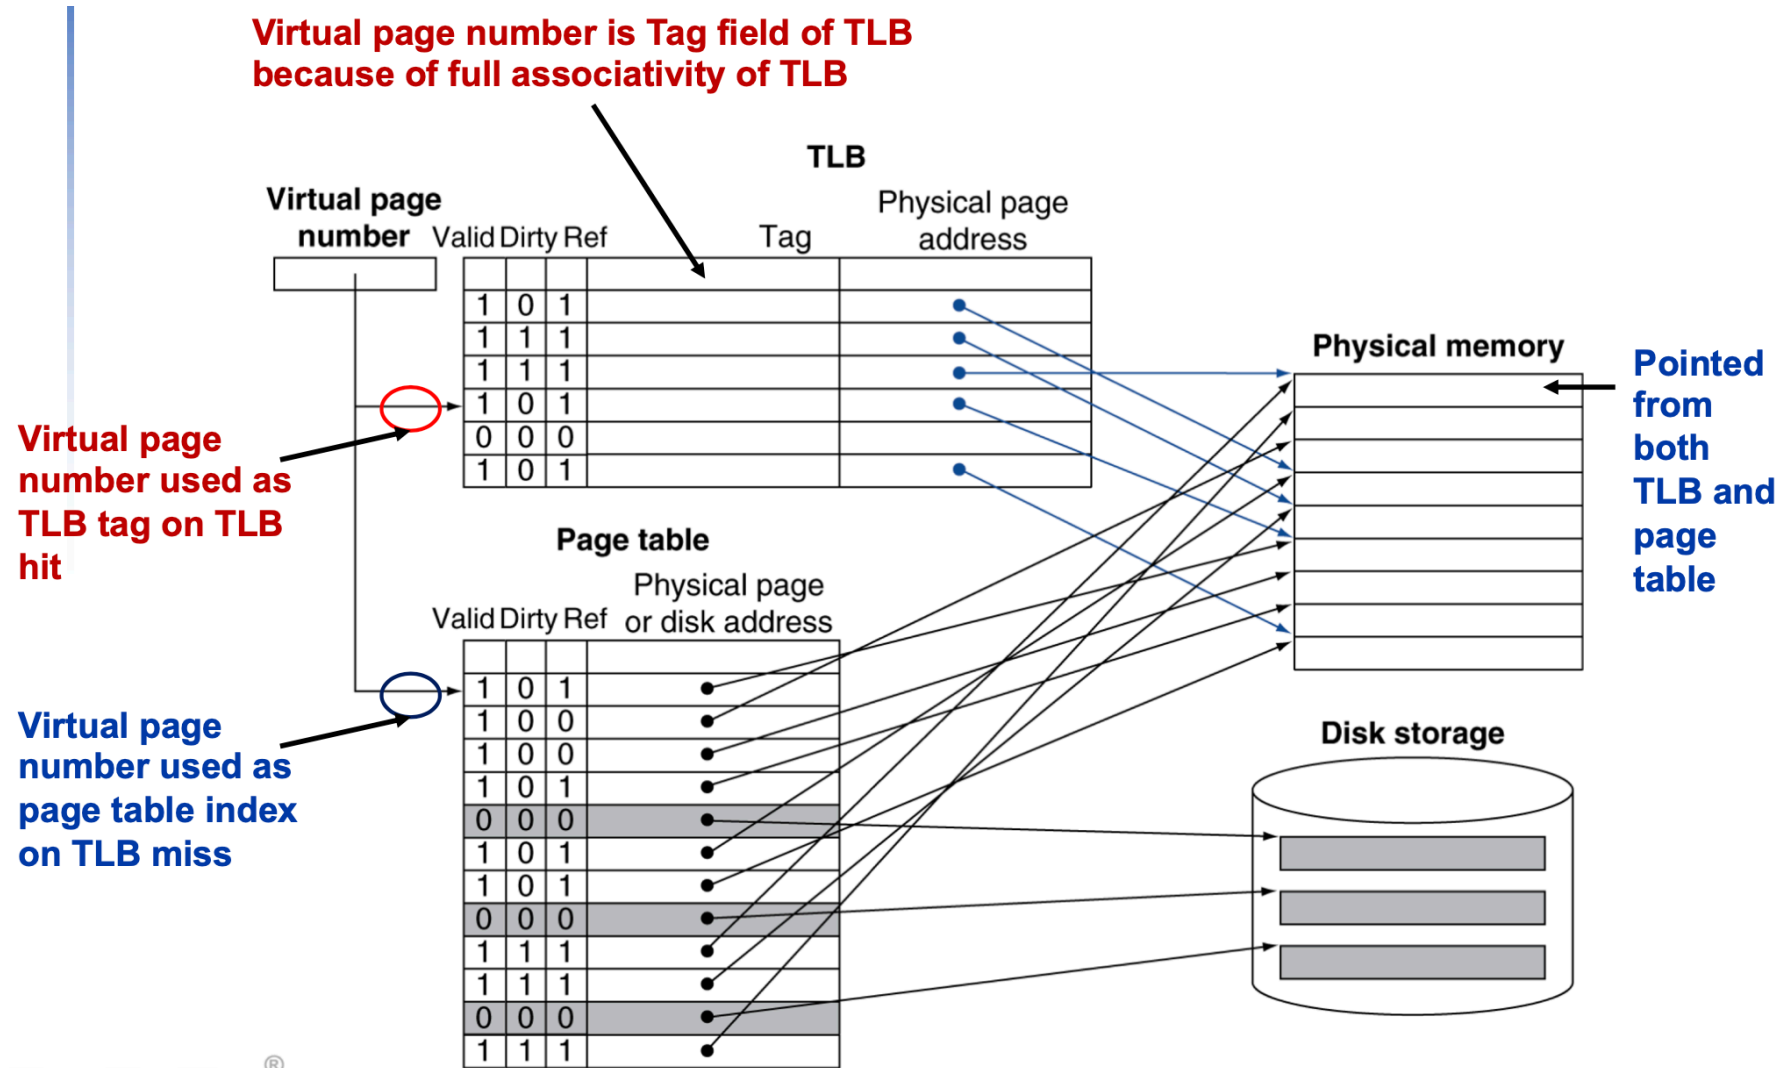
- Full name of TLB?
  - Translation look-aside buffer

- Matching:



# Key concept 3: TLB

## How to use TLB? (T14 – Page 24)



# Key concept 3: TLB

## TLB Hit & Miss

### TLB Hit

- Provide physical address
- Reference bit 0/1 → 1
- If write:
  - dirty bit 0/1 → 1
- Go to next instruction

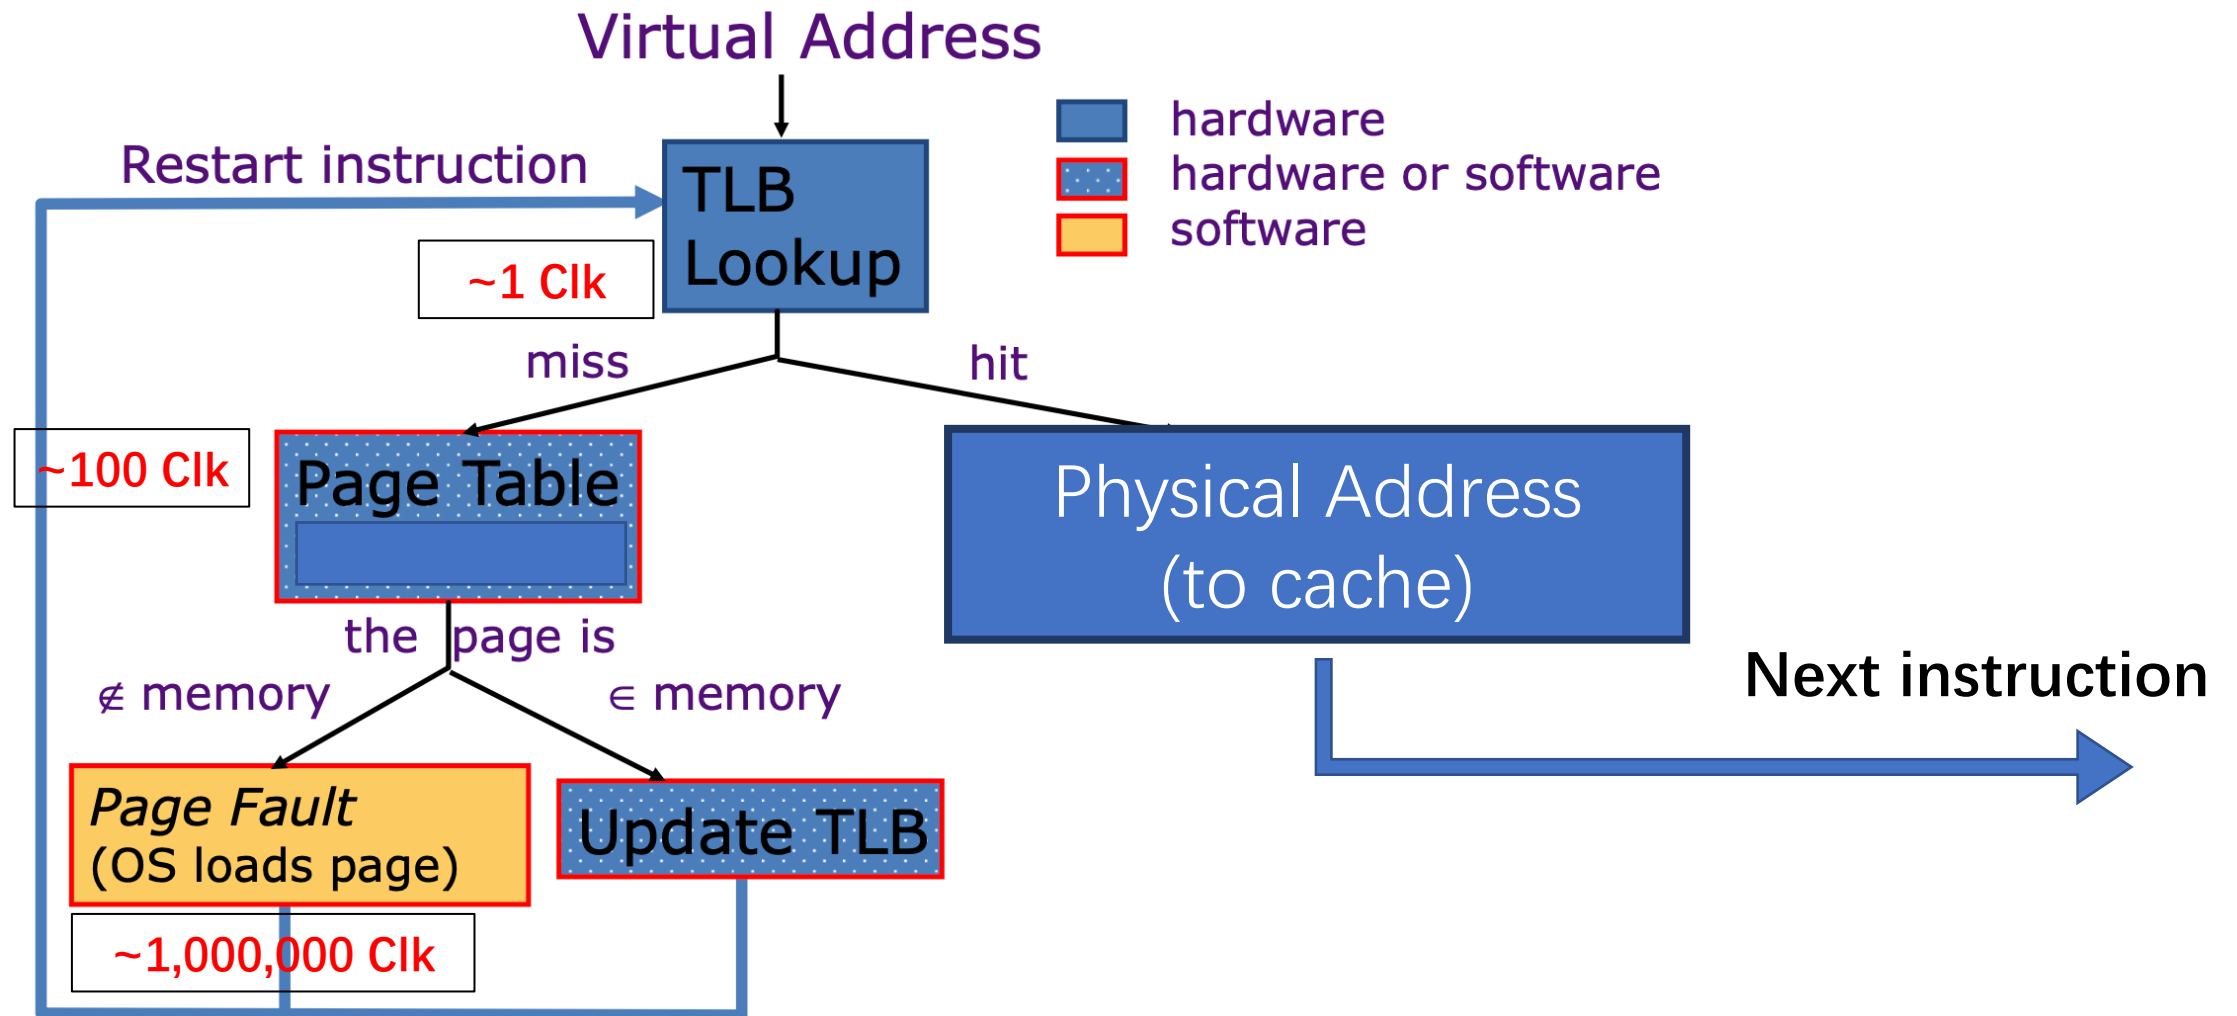
### TLB Miss

- Page table hit:
  - Load page table entry to TLB
- Page table miss:
  - Page fault exception
  - OS loads the page from disk
  - Update page table & TLB
- Restart current instruction



# Key concept 3: TLB

## TLB Hit & Miss



# Key concept 3: TLB

## Class exercise 2 – Homework 8.3.1

- 4 KB pages / 4-entry fully associative TLB / LRU replacement.
- If pages must be brought in from disk, increment to the next largest page number.
- Access to the memory by the following virtual address (decimal): 12948, 49419, 46814, 13975, 40004, 12707, 52236
- Convert to hexadecimal
- 0x3294, 0xC10B, 0xB6DE, 0x3697,
- 0x9C44, 0x31A3, 0xCC0C

V	Tag	PPN
1	11	12
1	7	4
1	3	6
0	4	9

TLB

#	V	PPN
0	1	5
1	0	Disk
2	0	Disk
3	1	6
4	1	9
5	1	1
6	0	Disk
7	1	4
8	0	Disk
9	0	Disk
10	1	3
11	1	12

Page Table

# Key concept 3: TLB

## Class exercise 2 – Homework 8.3.1

- 4 KB pages / 4-entry fully associative TLB / LRU replacement.
- **0x3294**, 0xC10B, 0xB6DE, 0x3697, 0x9C44, 0x31A3, 0xCC0C
- Physical address: 0x6294

TLB Hit

V	Tag	PPN
1	11	12
1	7	4
<b>1</b>	<b>3</b>	<b>6</b>
0	4	9

TLB

#	V	PPN
0	1	5
1	0	Disk
2	0	Disk
3	1	6
4	1	9
5	1	1
6	0	Disk
7	1	4
8	0	Disk
9	0	Disk
10	1	3
11	1	12

Page Table

# Key concept 3: TLB

## Class exercise 2 – Homework 8.3.1

- 4 KB pages / 4-entry fully associative TLB / LRU replacement.
- 0x3294, 0xC10B, 0xB6DE, 0x3697, 0x9C44, 0x31A3, 0xCC0C
- Physical address: 0xD10B

Page Fault

V	Tag	PPN
1	11	12
1	7	4
1	3	6
1	12	13

TLB

#	V	PPN
0	1	5
1	0	Disk
2	0	Disk
3	1	6
4	1	9
5	1	1
6	0	Disk
7	1	4
8	0	Disk
9	0	Disk
10	1	3
11	1	12
12	1	13

Page Table

# Key concept 3: TLB

## Class exercise 2 – Homework 8.3.1

- 4 KB pages / 4-entry fully associative TLB / LRU replacement.
- 0x3294, 0xC10B, 0xB6DE, 0x3697, 0x9C44, 0x31A3, 0xCC0C
- Physical address: 0xC6DE

TLB Hit

V	Tag	PPN
1	11	12
1	7	4
1	3	6
1	12	13

TLB

#	V	PPN
0	1	5
1	0	Disk
2	0	Disk
3	1	6
4	1	9
5	1	1
6	0	Disk
7	1	4
8	0	Disk
9	0	Disk
10	1	3
11	1	12
12	1	13

Page Table

# Key concept 3: TLB

## Class exercise 2 – Homework 8.3.1

- 4 KB pages / 4-entry fully associative TLB / LRU replacement.
- 0x3294, 0xC10B, 0xB6DE, 0x3697, 0x9C44, 0x31A3, 0xCC0C
- Physical address: 0x6697

TLB Hit

V	Tag	PPN
1	11	12
1	7	4
1	3	6
1	12	13

TLB

#	V	PPN
0	1	5
1	0	Disk
2	0	Disk
3	1	6
4	1	9
5	1	1
6	0	Disk
7	1	4
8	0	Disk
9	0	Disk
10	1	3
11	1	12
12	1	13

Page Table

# Key concept 3: TLB

## Class exercise 2 – Homework 8.3.1

- 4 KB pages / 4-entry fully associative TLB / LRU replacement.
- 0x3294, 0xC10B, 0xB6DE, 0x3697, 0x9C44, 0x31A3, 0xCC0C
- Physical address: 0xEC44

Page Fault

V	Tag	PPN
1	11	12
1	9	14
1	3	6
1	12	13

TLB

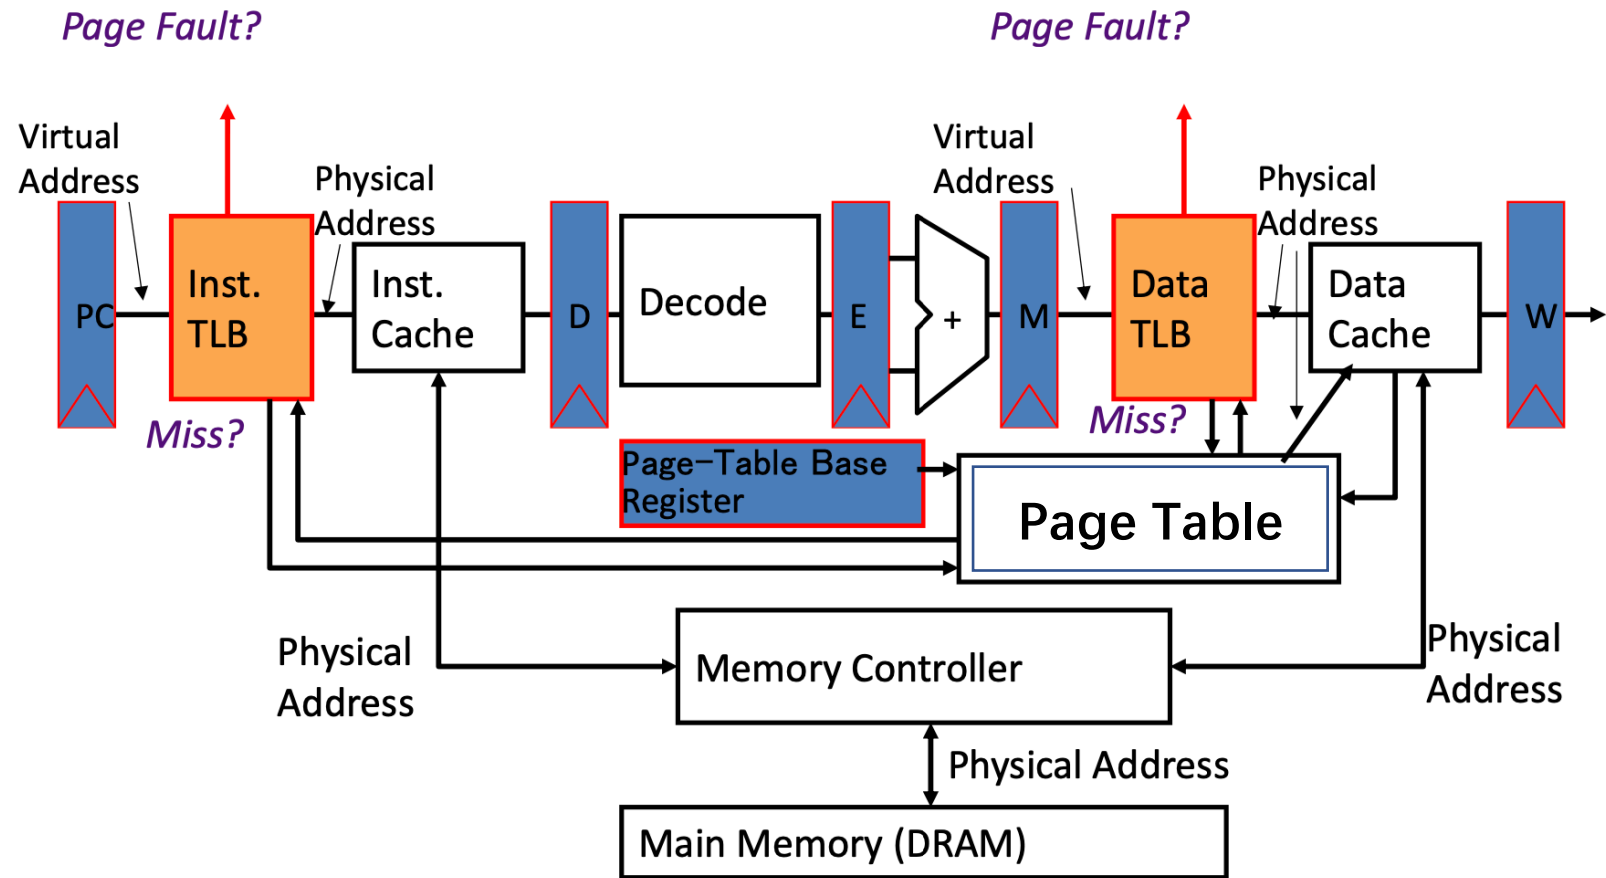
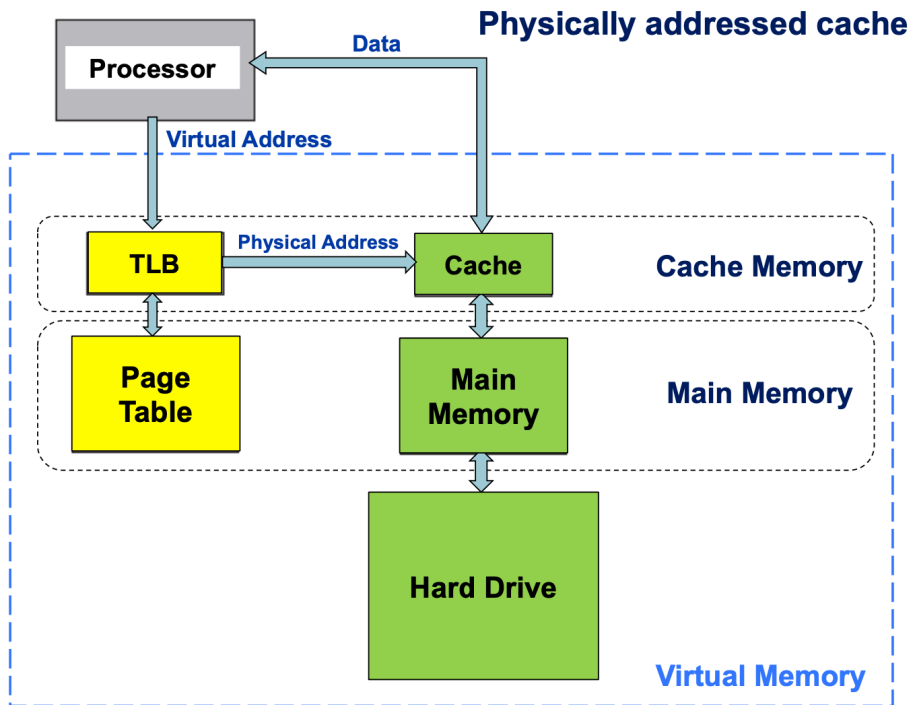
#	V	PPN
0	1	5
1	0	Disk
2	0	Disk
3	1	6
4	1	9
5	1	1
6	0	Disk
7	1	4
8	0	Disk
9	1	14
10	1	3
11	1	12
12	1	13

Page Table

The figure on the right will not be covered in the exam.

# Key concept 3: TLB

## The big picture (Design #1)



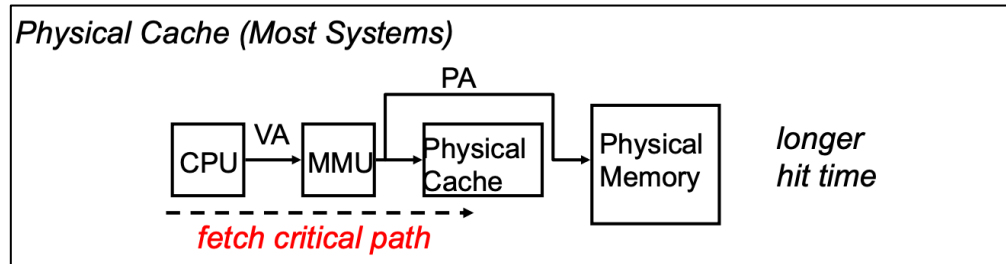


The figures on the left will not be covered in the exam.

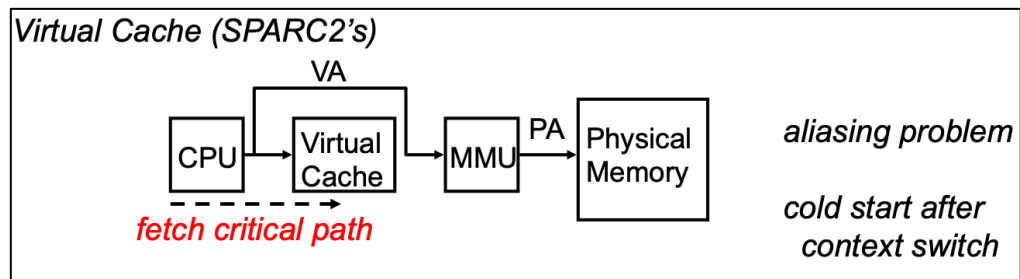
# Key concept 3: TLB

## The big picture (Design #2)

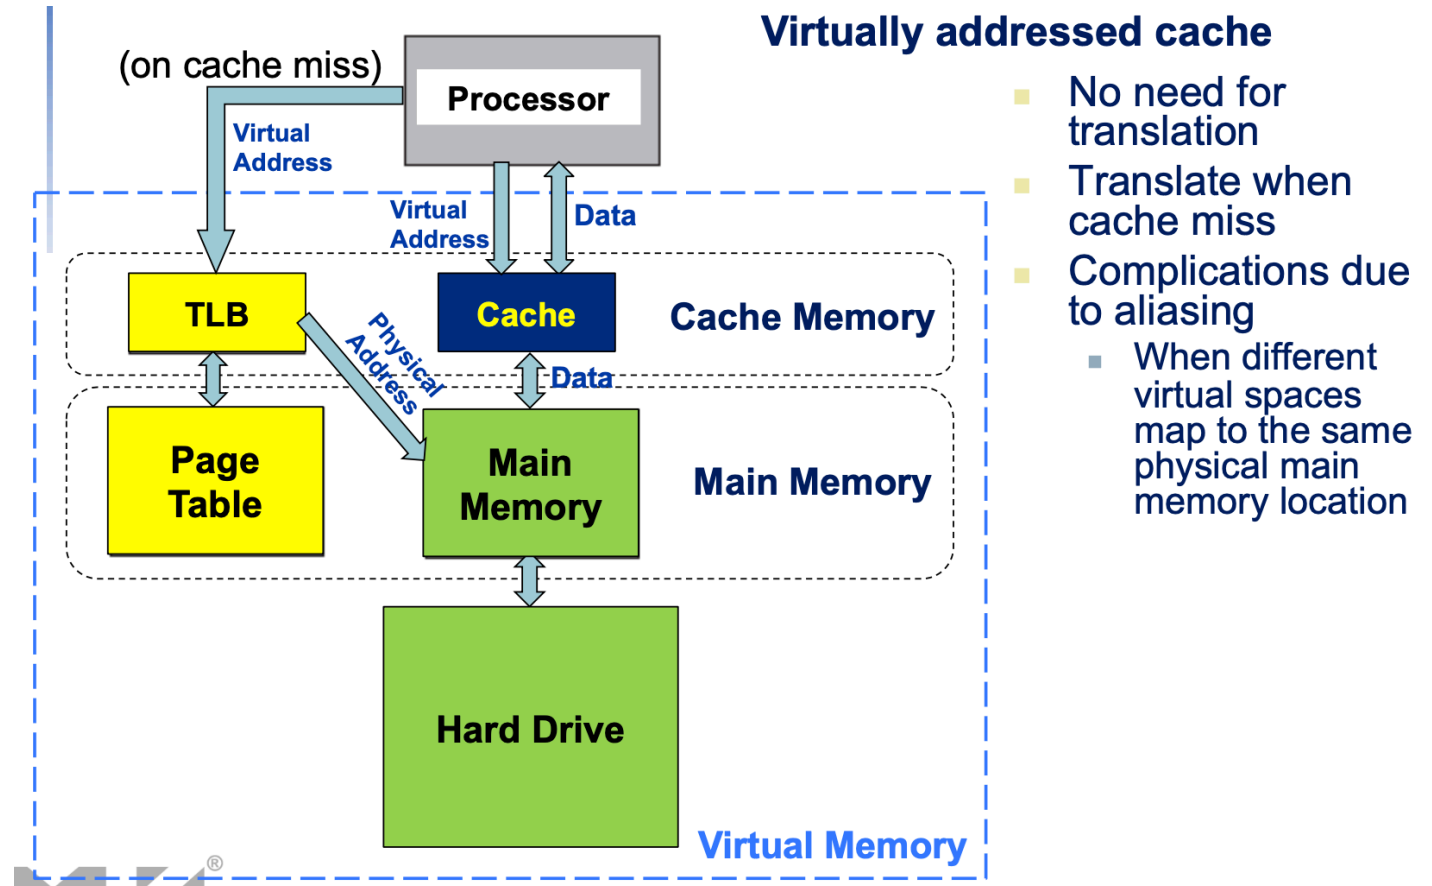
Get physical address in Design #1



Get physical address in Design #2



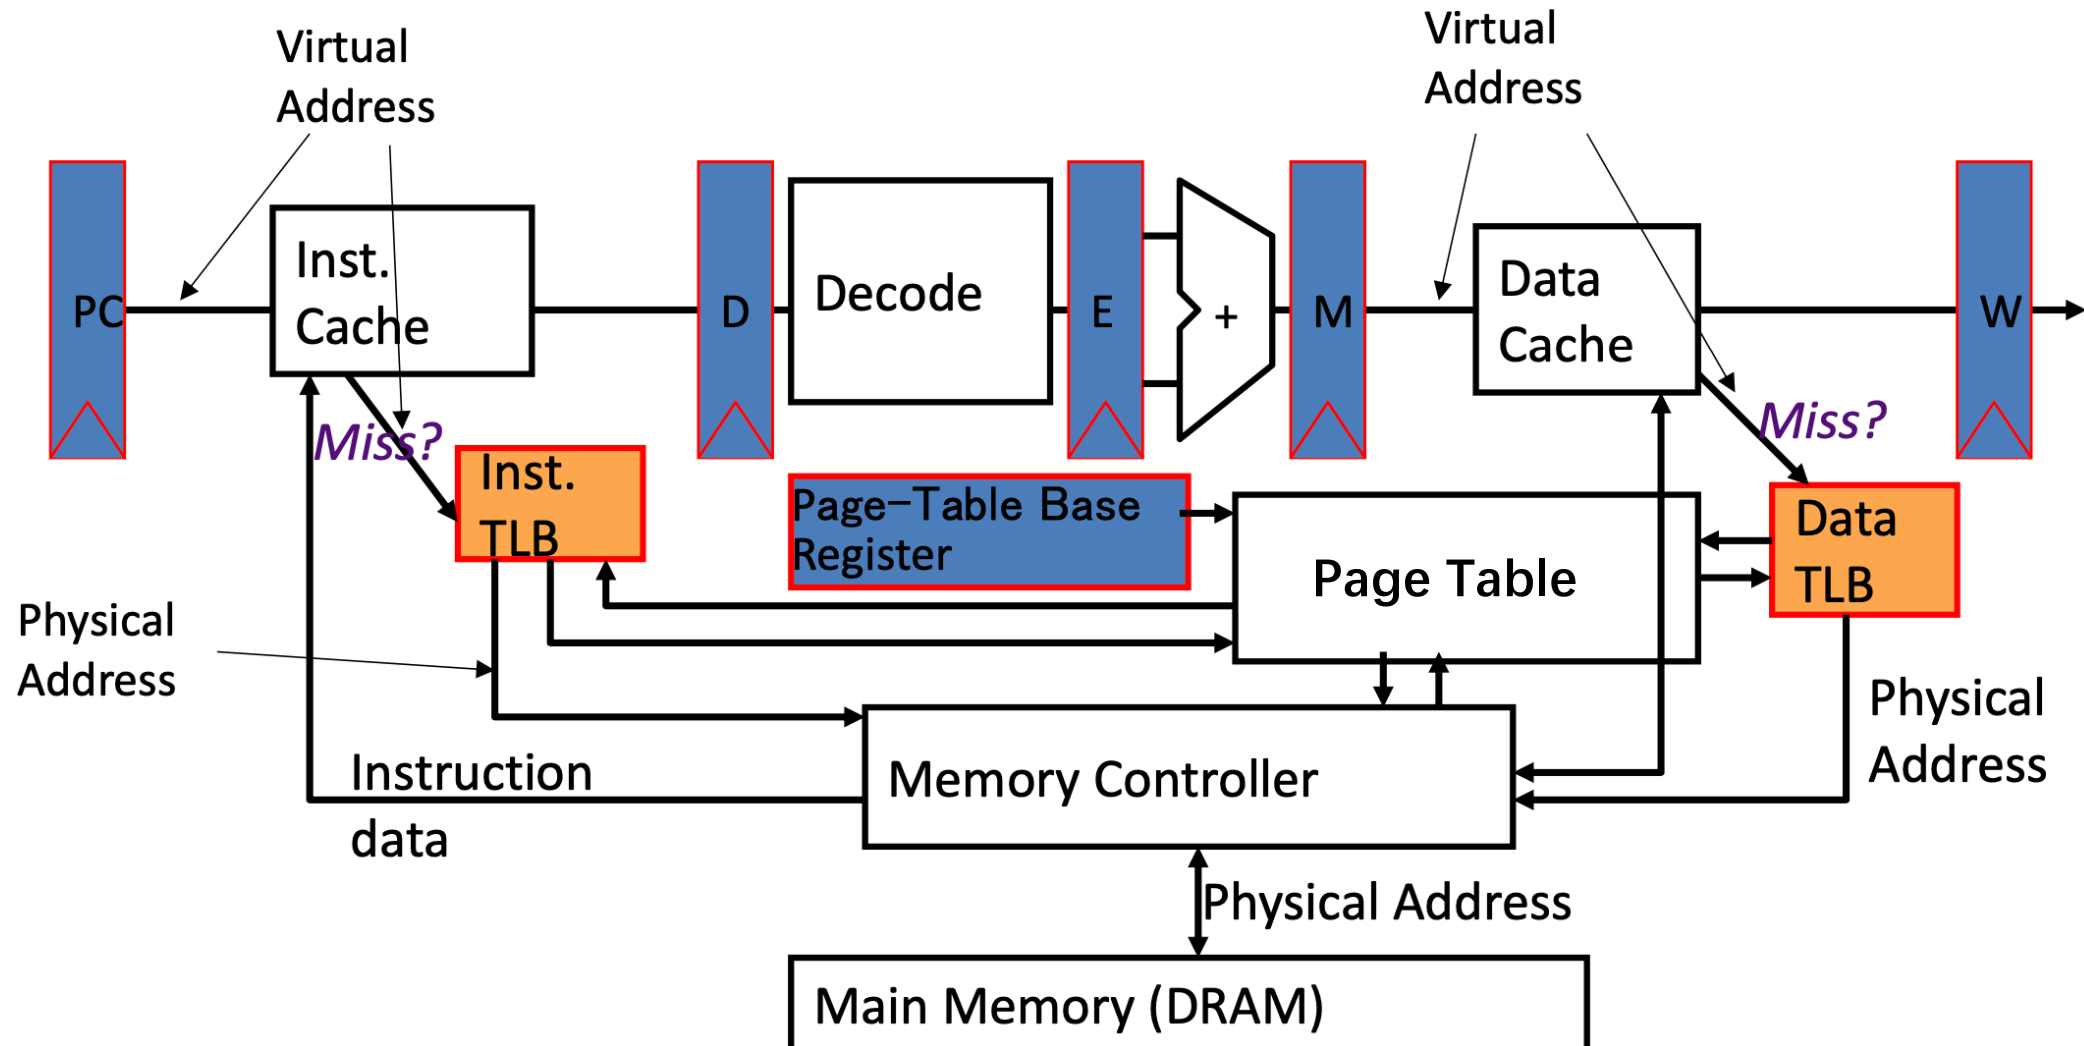
The big picture of design #2



This slide will not be covered in the exam.

# Key concept 3: TLB

## The big picture (Design #2)



# Thank you!

## Reference

- VE370 Lecture slides (T14 Virtual Memory).
- Princeton ELE475 Lecture slides (L10 Address Translation and Protection).
- UM EECS470 Lecture slides (L20 Virtual Memory).
- Caches. Accessed: 2020-11-12. <https://uops.info/cache.html>